

Lecture 5

Scalable PCA

Dimensionality Reduction & Factor Analysis

Haiping Lu

<https://github.com/haipinglu/ScalableML>

COM6012 Scalable Machine Learning

Spring 2019


Week 5 Contents / Objectives

- PCA - Dimensionality Reduction
- SVD – Factor Analysis
- Scalable PCA in Spark
- More on Scala

Week 5 Contents / Objectives

- **PCA - Dimensionality Reduction**
- SVD – Factor Analysis
- Scalable PCA in Spark
- More on Scala

Dimensionality Reduction



- **Assumption:** Data lies on or near a low d -dimensional subspace
- Axes of this subspace are effective representation of the data

Why Reduce Dimensions?

Why reduce dimensions?

- Discover hidden correlations/topics
 - Words that occur commonly together
- Remove redundant and noisy features
 - Not all words are useful
- Interpretation and visualization
- Easier storage and processing of the data

Dimensionality Reduction

- Raw data is complex and high-dimensional
- Dimensionality reduction describes the data using a simpler, more compact representation
- This representation may make interesting patterns in the data clearer or easier to see

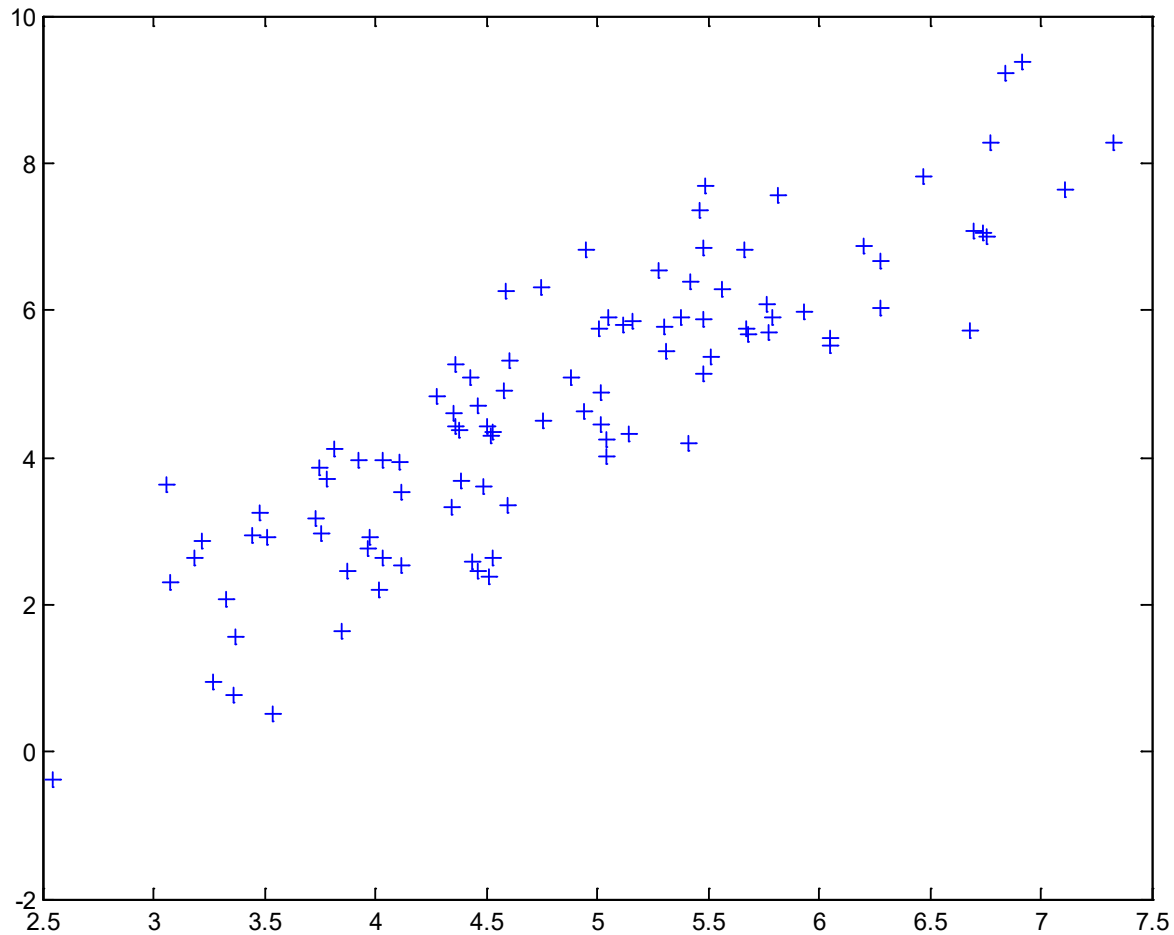
Dimensionality Reduction

- Goal: Find a ‘better’ representation for data
- How do we define ‘better’?
- For example
 - Minimise reconstruction error
 - Maximise variance
 - **They give the same solution → PCA!**

PCA Algorithm

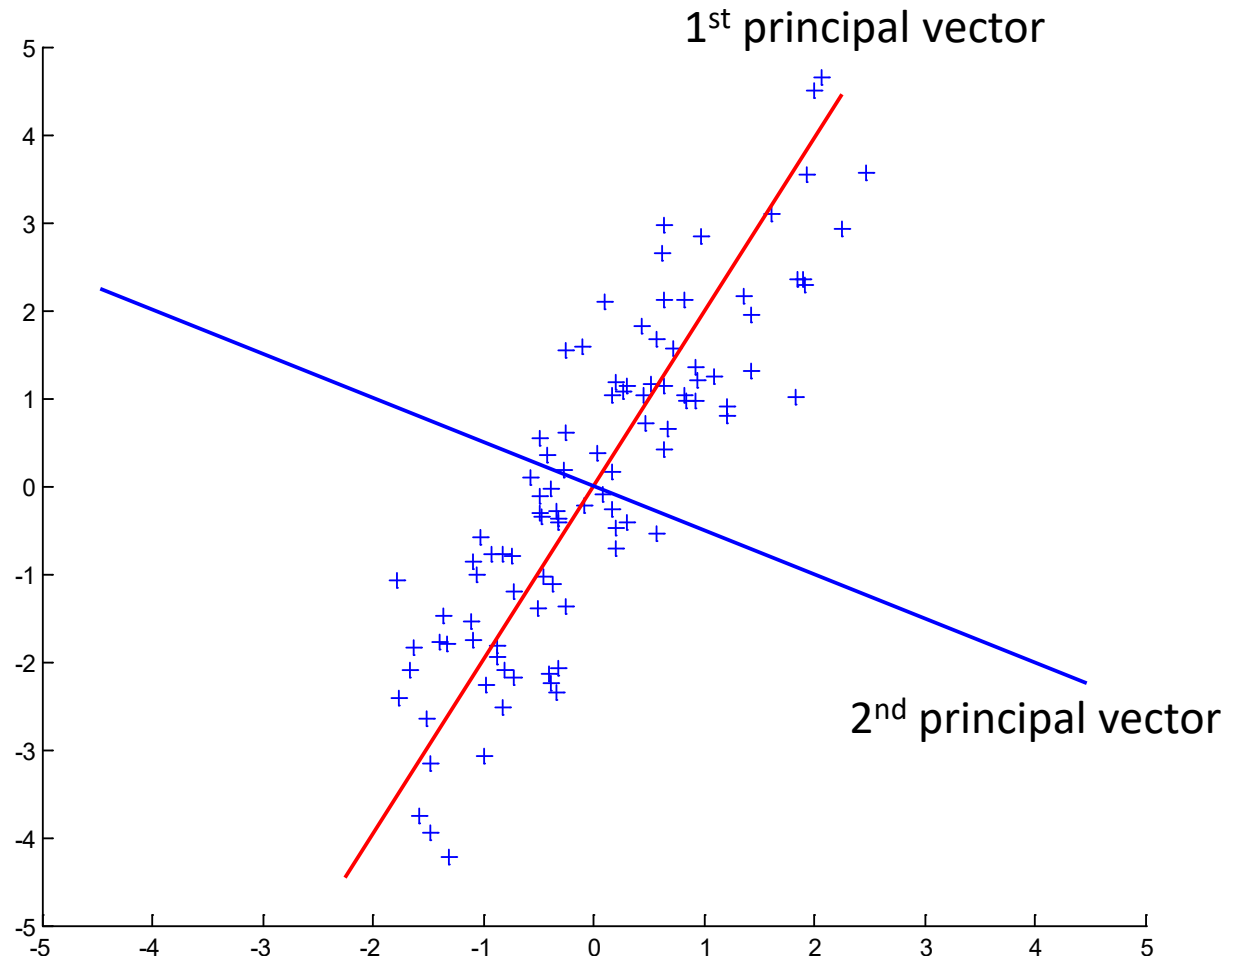
- Input: N data points, each $\rightarrow D$ -dimensional vector
- PCA algorithm
 - 1. $\mathbf{X}_0 \leftarrow$ Form $N \times D$ data matrix, with one row vector \mathbf{x}_n per data point
 - 2. \mathbf{X} : subtract mean \mathbf{x} from each row vector \mathbf{x}_n in \mathbf{X}_0
 - 3. $\Sigma \leftarrow \mathbf{X}^T \mathbf{X}$ Gramian (scatter) matrix for \mathbf{X}
 - Find eigenvectors and eigenvalues of Σ
 - PCs $\mathbf{U} (D \times d) \leftarrow$ the d eigenvectors with largest eigenvalues
- PCA feature for \mathbf{y} D -dim: $\mathbf{U}^T \mathbf{y}$ (d -dimensional)
 - Zero correlations, ordered by variance

2D Data



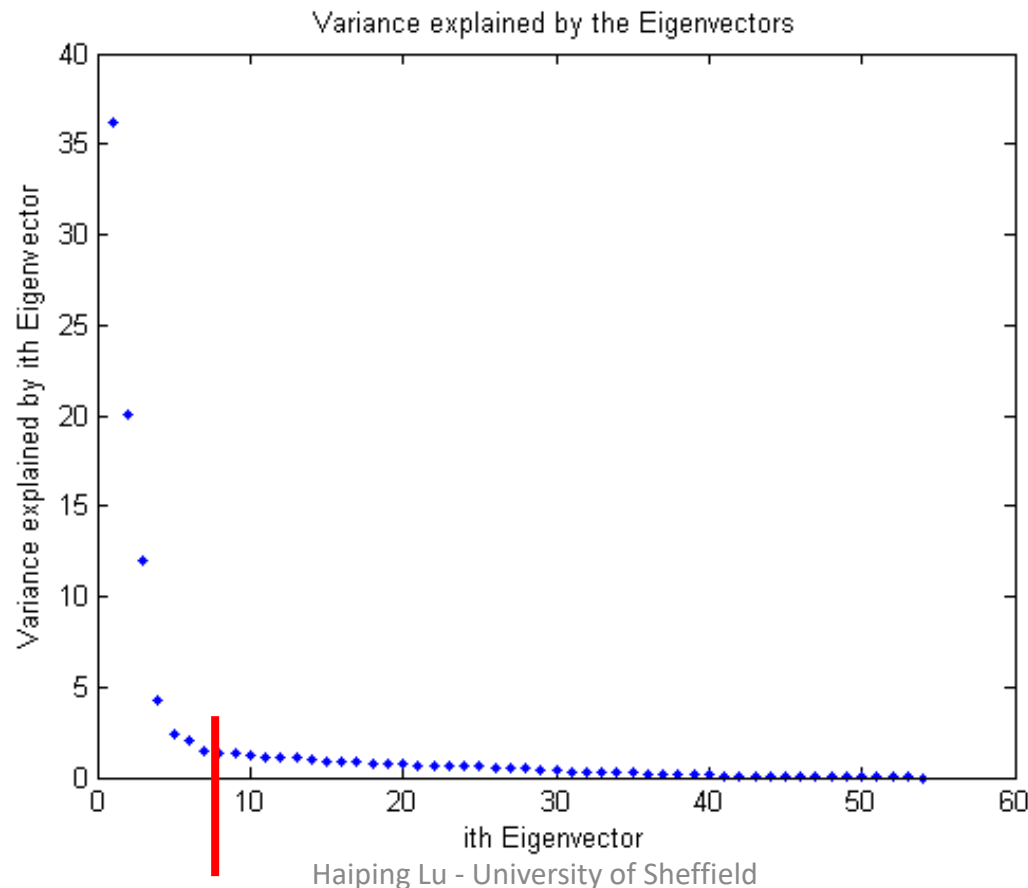
Principal Components

- The best axis to project
- Minimum RMS error
- Principal vectors are **orthogonal**



How Many Components?

- Check the distribution of eigen-values
- Take enough many eigen-vectors to cover 80-90% of the variance



Other Practical Tips

- PCA assumptions (linearity, orthogonality) not always appropriate
- Various extensions to PCA with different underlying assumptions, e.g., manifold learning, Kernel PCA, ICA
- Centring is crucial, i.e., we must preprocess data so that all features have zero mean before applying PCA
- PCA results dependent on scaling of data
- Data is sometimes rescaled in practice before applying PCA

Problems and Limitations

- What if very large dimensional data?
 - e.g., Images ($D \geq 10^4 = 100 \times 100$)
- Problem:
 - Gramian matrix Σ is size (D^2)
 - $D=10^4 \rightarrow |\Sigma| = 10^8$
- Singular Value Decomposition (SVD)!
 - Efficient algorithms available
 - Some implementations find just top d eigenvectors

Week 5 Contents / Objectives

- PCA - Dimensionality Reduction
- **SVD – Factor Analysis**
- Scalable PCA in Spark
- More on Scala

Singular Value Decomposition

- Factorization (decomposition) problem
 - #1: Find concepts/topics/genres → Factor Analysis
 - #2: Reduce dimensionality

term	data	information	retrieval	brain	lung
document					
CS-TR1	1	1	1	0	0
CS-TR2	2	2	2	0	0
CS-TR3	1	1	1	0	0
CS-TR4	5	5	5	0	0
MED-TR1	0	0	0	2	2
MED-TR2	0	0	0	3	3
MED-TR3	0	0	0	1	1

The above matrix is actually “2-dimensional.” All rows can be reconstructed by scaling $[1 \ 1 \ 1 \ 0 \ 0]$ or $[0 \ 0 \ 0 \ 1 \ 1]$: $D=5 \rightarrow d=2$

SVD - Definition

$$\mathbf{A}_{[n \times m]} = \mathbf{U}_{[n \times r]} \mathbf{\Lambda}_{[r \times r]} (\mathbf{V}_{[m \times r]})^T$$

- \mathbf{A} : $n \times m$ matrix (e.g., n documents, m terms)
- \mathbf{U} : $n \times r$ matrix (n documents, r concepts)
- $\mathbf{\Lambda}$: $r \times r$ diagonal matrix (strength of each ‘concept’) (r : rank of the matrix)
- \mathbf{V} : $m \times r$ matrix (m terms, r concepts)

SVD - Properties

Always possible to decompose matrix \mathbf{A} into $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$,
where

- $\mathbf{U}, \mathbf{\Lambda}, \mathbf{V}$: unique (*)
- \mathbf{U}, \mathbf{V} : column orthonormal (i.e., columns are unit vectors, orthogonal to each other)
 - $\mathbf{U}^T \mathbf{U} = \mathbf{I}; \mathbf{V}^T \mathbf{V} = \mathbf{I}$ (\mathbf{I} : identity matrix)
- $\mathbf{\Lambda}$: singular value are positive, and sorted in decreasing order

SVD \leftrightarrow Eigen-decomposition

- SVD gives us:
 - $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$
- Eigen-decomposition:
 - $\mathbf{B} = \mathbf{W} \mathbf{\Sigma} \mathbf{W}^T$
 - $\mathbf{U}, \mathbf{V}, \mathbf{W}$ are orthonormal ($\mathbf{U}^T \mathbf{U} = \mathbf{I}$),
 - $\mathbf{\Lambda}, \mathbf{\Sigma}$ are diagonal
- Relationship:
 - $\mathbf{A} \mathbf{A}^T = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T (\mathbf{U} \mathbf{\Lambda} \mathbf{V}^T)^T = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T (\mathbf{V} \mathbf{\Lambda}^T \mathbf{U}^T) = \mathbf{U} \mathbf{\Lambda} \mathbf{\Lambda}^T \mathbf{U}^T$
 - $\mathbf{A}^T \mathbf{A} = \mathbf{V} \mathbf{\Lambda}^T \mathbf{U}^T (\mathbf{U} \mathbf{\Lambda} \mathbf{V}^T) = \mathbf{V} \mathbf{\Lambda} \mathbf{\Lambda}^T \mathbf{V}^T = \mathbf{V} \mathbf{\Lambda}^2 \mathbf{V}^T$
 - $\mathbf{B} = \mathbf{A}^T \mathbf{A} = \mathbf{W} \mathbf{\Sigma} \mathbf{W}^T$

SVD for PCA

- PCA by SVD:
 - 1. $\mathbf{X}_0 \leftarrow$ Form $N \times d$ data matrix, with one row vector \mathbf{x}_n per data point
 - 2. \mathbf{X} subtract mean \mathbf{x} from each row vector \mathbf{x}_n in \mathbf{X}_0
 - 3. $\mathbf{U} \mathbf{\Lambda} \mathbf{V}^T \leftarrow$ SVD of \mathbf{X}
 - The right singular vectors \mathbf{V} of \mathbf{X} are equivalent to the eigenvectors of $\mathbf{X}^T \mathbf{X} \rightarrow$ the PCs
 - The singular values in $\mathbf{\Lambda}$ are equal to the square roots of the eigenvalues of $\mathbf{X}^T \mathbf{X}$

SVD - Properties

‘spectral decomposition’ of the matrix:

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} | & | \\ u_1 & u_2 \\ | & | \end{bmatrix} \times \begin{bmatrix} \lambda_1 & \emptyset \\ \emptyset & \lambda_2 \end{bmatrix} \times \begin{bmatrix} \text{---} v_1 \text{---} \\ \text{---} v_2 \text{---} \end{bmatrix}$$

SVD - Interpretation

‘documents’, ‘terms’ and ‘concepts’:

- U : document-to-concept similarity matrix
- V : term-to-concept similarity matrix
- Λ : its diagonal elements: ‘strength’ of each concept

Projection:

- Best axis to project on: (‘best’ = min sum of squares of projection errors)

SVD - Example

- $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$ - example:

term	data	information	retrieval	brain	lung
document					
CS-TR1	1	1	1	0	0
CS-TR2	2	2	2	0	0
CS-TR3	1	1	1	0	0
CS-TR4	5	5	5	0	0
MED-TR1	0	0	0	2	2
MED-TR2	0	0	0	3	3
MED-TR3	0	0	0	1	1

$$\begin{array}{c}
 \uparrow \\
 \text{CS} \\
 \downarrow \\
 \uparrow \\
 \text{MD} \\
 \downarrow
 \end{array}
 \begin{bmatrix}
 1 & 1 & 1 & 0 & 0 \\
 2 & 2 & 2 & 0 & 0 \\
 1 & 1 & 1 & 0 & 0 \\
 5 & 5 & 5 & 0 & 0 \\
 0 & 0 & 0 & 2 & 2 \\
 0 & 0 & 0 & 3 & 3 \\
 0 & 0 & 0 & 1 & 1
 \end{bmatrix}
 =
 \begin{bmatrix}
 0.18 & 0 \\
 0.36 & 0 \\
 0.18 & 0 \\
 0.90 & 0 \\
 0 & 0.53 \\
 0 & 0.80 \\
 0 & 0.27
 \end{bmatrix}
 \times
 \begin{bmatrix}
 9.64 & 0 \\
 0 & 5.29
 \end{bmatrix}
 \times
 \begin{bmatrix}
 0.58 & 0.58 & 0.58 & 0 & 0 \\
 0 & 0 & 0 & 0.71 & 0.71
 \end{bmatrix}$$

SVD - Example

- $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$ - example:

doc-to-concept
similarity matrix

CS-concept MD-concept

↑ CS
↓
↑ MD
↓

data inf. ↓ retrieval brain lung

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix} \times \begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix} \times \begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$$

SVD - Example

- $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$ - example:

retrieval
inf. ↓ brain lung

data

‘strength’ of CS-concept

↑

CS

↓

↑

MD

↓

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix}$$
 $=$

$$\begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix}$$
 \times

$$\begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix}$$
 \times

$$\begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$$

SVD - Example

- $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$ - example:

retrieval

inf. ↓

data brain lung

↑ CS

↓

↑ MD

↓

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix} \times \begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix} \times \begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$$

term-to-concept
similarity matrix

CS-concept

SVD – Dimensionality Reduction

- Q: how exactly is (**further**) dim. reduction done?
- A: set the smallest singular values to zero:
- Note: **3 zero singular values** already removed

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix} \times \begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix} \times \begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$$

Diagram illustrating SVD decomposition for dimensionality reduction. The original matrix is decomposed into three matrices: a 7x2 matrix of left singular vectors, a 2x2 diagonal matrix of singular values, and a 2x5 matrix of right singular vectors. The singular values 9.64 and 5.29 are shown, with the second singular value (5.29) being crossed out with a blue 'X', indicating it is the smallest and thus the target for reduction. The right singular vectors matrix shows the first two columns (0.58, 0.58, 0.58, 0, 0) and the last two columns (0, 0, 0, 0.71, 0.71) crossed out with a red 'X', indicating they are the smallest and thus the target for reduction.

SVD - Dimensionality Reduction

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} \sim \begin{bmatrix} 0.18 \\ 0.36 \\ 0.18 \\ 0.90 \\ 0 \\ 0 \\ 0 \end{bmatrix} \times \begin{bmatrix} 9.64 \end{bmatrix} \times \begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \end{bmatrix}$$

SVD - Dimensionality Reduction

- Best rank-1 approximation

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} \sim \begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Week 5 Contents / Objectives

- PCA - Dimensionality Reduction
- SVD – Factor Analysis
- **Scalable PCA in Spark**
- More on Scala

PCA & SVD in Spark MLlib

- Not scalable: `computePrincipalComponents()` from `RowMatrix`
- **Scalable:** `computeSVD()` from `RowMatrix`
- Code:
<https://github.com/apache/spark/blob/v2.3.2/mllib/src/main/scala/org/apache/spark/mllib/linalg/distributed/RowMatrix.scala>
- Documentation:
<https://spark.apache.org/docs/2.3.2/api/scala/index.html#org.apache.spark.mllib.linalg.distributed.RowMatrix>

PCA in Spark MLlib (RDD)

- <https://spark.apache.org/docs/2.3.2/mllib-dimensionality-reduction.html>

```
val mat: RowMatrix = new RowMatrix(dataRDD)

// Compute the top 4 principal components.
// Principal components are stored in a local dense matrix.
val pc: Matrix = mat.computePrincipalComponents(4)
```

- Not scalable, local computation

```
val brzSvd.SVD(u: BDM[Double], s: BDV[Double], _) = brzSvd(Cov)
```

PCA in Spark ML (DF)

- Now in

<https://spark.apache.org/docs/2.3.2/ml-features.html#pca>

- Under features
- Not scalable

```
val pca = new PCA()  
  .setInputCol("features")  
  .setOutputCol("pcaFeatures")  
  .setK(3)  
  .fit(df)
```


SVD in Spark MLlib (RDD)

- <https://spark.apache.org/docs/2.3.2/mllib-dimensionality-reduction.html>
- With distributed implementations

```
val mat: RowMatrix = new RowMatrix(dataRDD)

// Compute the top 5 singular values and corresponding singular vectors.
val svd: SingularValueDecomposition[RowMatrix, Matrix] = mat.computeSVD(5, computeU = true)
val U: RowMatrix = svd.U // The U factor is a RowMatrix.
val s: Vector = svd.s // The singular values are stored in a local dense vector.
val V: Matrix = svd.V // The V factor is a local dense matrix.
```

SVD in Spark MLlib (RDD)

- An $m \times n$ data matrix \mathbf{A} with $m > n$ (note different notations)
- For large matrices, usually we don't need the complete factorization but only the top k singular values and its associated singular vectors.
- Save storage, de-noise and recover the low-rank structure of the matrix (dimensionality reduction)

SVD in Spark MLlib (RDD)

- An $m \times n$ data matrix \mathbf{A}
- Assume $m > n$, SVD $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$
- The singular values and the right singular vectors are derived from the eigenvalues and the eigenvectors of $\mathbf{A}^T \mathbf{A}$ (which is smaller than \mathbf{A})
- The left singular vectors are computed via matrix multiplication as $\mathbf{U} = \mathbf{A} \mathbf{V} \mathbf{\Lambda}^{-1}$, if requested by the user via the computeU parameter

Selection of SVD Computation

- Auto
- If n is small ($n < 100$) or k is large compared with n ($k > n/2$)
 - Compute $\mathbf{A}^T \mathbf{A}$ first and then compute its top eigenvalues and eigenvectors **locally** on the driver
- Otherwise
 - Compute $\mathbf{A}^T \mathbf{A} \mathbf{v}$ in a distributive way and send it to ARPACK to compute the top eigenvalues and eigenvectors on the driver node

Selection of SVD Computation

- Auto (default)

```
if (n < 100 || (k > n / 2 && n <= 15000)) {  
    // If n is small or k is large compared with n, we better compute the Gramian matrix first  
    // and then compute its eigenvalues locally, instead of making multiple passes.  
    if (k < n / 3) {  
        SVDMode.LocalARPACK  
    } else {  
        SVDMode.LocalLAPACK  
    }  
} else {  
    // If k is small compared with n, we use ARPACK with distributed multiplication.  
    SVDMode.DistARPACK  
}
```

Selection of SVD Computation

- Specify computeMode (private)

```
case "local-svd" => SVDMODE.LocalLAPACK  
case "local-eigs" => SVDMODE.LocalARPACK  
case "dist-eigs" => SVDMODE.DistARPACK
```

Selection of SVD Computation

- computeMode (note brzSvd.SVD is local)

```
// Compute the eigen-decomposition of A' * A.
val (sigmaSquares: BDV[Double], u: BDM[Double]) = computeMode match {
  case SVDMode.LocalARPACK =>
    require(k < n, s"k must be smaller than n in local-eigs mode but got k=$k and n=$n.")
    val G = computeGramianMatrix().asBreeze.asInstanceOf[BDM[Double]]
    EigenValueDecomposition.symmetricEigs(v => G * v, n, k, tol, maxIter)
  case SVDMode.LocalLAPACK =>
    // breeze (v0.10) svd latent constraint, 7 * n * n + 4 * n < Int.MaxValue
    require(n < 17515, s"$n exceeds the breeze svd capability")
    val G = computeGramianMatrix().asBreeze.asInstanceOf[BDM[Double]]
    val brzSvd.SVD(uFull: BDM[Double], sigmaSquaresFull: BDV[Double], _) = brzSvd(G)
    (sigmaSquaresFull, uFull)
  case SVDMode.DistARPACK =>
    if (rows.getStorageLevel == StorageLevel.NONE) {
      logWarning("The input data is not directly cached, which may hurt performance if its"
        + " parent RDDs are also uncached.")
    }
    require(k < n, s"k must be smaller than n in dist-eigs mode but got k=$k and n=$n.")
    EigenValueDecomposition.symmetricEigs(multiplyGramianMatrixBy, n, k, tol, maxIter)
}
```

Week 5 Contents / Objectives

- PCA - Dimensionality Reduction
- SVD – Factor Analysis
- Scalable PCA in Spark
- **More on Scala**

Scala (Scalable language)

- A pure object-oriented language. Conceptually, every value is an object and every operation is a method-call.
- A functional language. Supports functions, immutable data structures and preference for immutability over mutation
- Seamlessly integrated with Java
 - Mixed Scala/Java projects
 - Use existing Java libraries
 - Use existing Java tools

Scala Basic Syntax

- When considering a Scala program, it can be defined as a collection of objects that communicate via invoking each other's methods.
- **Object** – same as in Java
- **Class** – same as in Java
- **Methods** – same as in Java
- **Fields** – Each object has its unique set of instant variables, which are called fields. An object's state is created by the values assigned to these fields.
- **Traits** – Like Java Interface. A trait encapsulates method and field definitions, which can then be reused by mixing them into classes.
- **Closure** – A **closure** is a function, whose return value depends on the value of one or more variables declared outside this function.

closure = function + environment

Scala is Statically Typed

- You don't have to specify a type in most cases
- Type Inference

```
val sum = 1 + 2 + 3
```

```
val nums = List(1, 2, 3)
```

```
val map = Map("abc" -> List(1,2,3))
```

Explicit Types

```
val sum: Int = 1 + 2 + 3
```

```
val nums: List[Int] = List(1, 2, 3)
```

```
val map: Map[String, List[Int]] = ...
```

Scala is High level

// Java - Check if string has uppercase character

```
boolean hasUpperCase = false;
```

```
for(int i = 0; i < name.length(); i++) {
```

```
    if(Character.isUpperCase(name.charAt(i))) {
```

```
        hasUpperCase = true;
```

```
        break;
```

```
    }
```

```
}
```

// Scala

```
val hasUpperCase = name.exists(_.isUpperCase)
```

Scala is Concise

// Java

```
public class Person {  
    private String name;  
    private int age;  
    public Person(String name, Int age) {  
        this.name = name;  
        this.age = age;  
    }  
    public String getName() {                // name getter  
        return name;  
    }  
    public int getAge() {                    // age getter  
        return age;  
    }  
    public void setName(String name) {      // name setter  
        this.name = name;  
    }  
    public void setAge(int age) {           // age setter  
        this.age = age;  
    }  
}
```

// Scala

```
class Person(var name: String, private var _age: Int) {  
    def age = _age                // Getter for age  
    def age_=(newAge: Int) {      // Setter for age  
        println("Changing age to: "+newAge)  
        _age = newAge  
    }  
}
```

Variables and Values

- **V**ariables: values stored can be changed

```
var foo = "foo"  
foo = "bar"    // okay
```

- Values: immutable variable

```
val foo = "foo"  
foo = "bar"    // nope
```

Scala is Functional

- First Class Functions. Functions are treated like objects:
 - passing functions as arguments to other functions
 - returning functions as the values from other functions
 - assigning functions to variables or storing them in data structures

```
// Lightweight anonymous functions
```

```
(x:Int) => x + 1
```

```
// Calling the anonymous function
```

```
val plusOne = (x:Int) => x + 1
```

```
plusOne(5) → 6
```

Scala is Functional

- Closures: a function whose return value depends on the value of one or more variables declared outside this function.

// plusFoo can reference any **values/variables** in scope

```
var foo = 1
```

```
val plusFoo = (x:Int) => x + foo
```

```
plusFoo(5)  → 6
```

// Changing foo changes the return value of plusFoo

```
foo = 5
```

```
plusFoo(5)  → 10
```


Scala is Functional

- Higher Order Functions
 - A function that does at least one of the following:
 - takes one or more functions as arguments
 - returns a function as its result

```
val plusOne = (x:Int) => x + 1
```

```
val nums = List(1,2,3)
```

```
// map takes a function: Int => T
```

```
nums.map(plusOne)           → List(2,3,4)
```

```
// Inline Anonymous
```

```
nums.map(x => x + 1)        → List(2,3,4)
```

```
// Short form
```

```
nums.map(_ + 1)            → List(2,3,4)
```

More Examples on Higher Order Functions

```
val nums = List(1,2,3,4)
```

```
// A few more examples for List class
```

```
nums.exists(_ == 2)           → true
```

```
nums.find(_ == 2)             → Some(2)
```

```
nums.indexOf(_ == 2)          → 1
```

```
// functions as parameters, apply f to the  
value "1"
```

```
def call(f: Int => Int) = f(1)
```

```
call(plusOne)                 → 2
```

```
call(x => x + 1)              → 2
```

```
call(_ + 1)                   → 2
```

The Usage of “_” in Scala

- In anonymous functions, the “_” acts as a placeholder for parameters

nums.map(x => x + 1) is equivalent to:

nums.map(_ + 1)

List(1,2,3,4,5).foreach(print(_)) is equivalent to:

List(1,2,3,4,5).foreach(a => print(a))

- You can use two or more underscores to refer different parameters.

val sum = List(1,2,3,4,5).reduceLeft(_ + _) is equivalent to:

val sum = List(1,2,3,4,5).reduceLeft((a, b) => a + b)

- The reduceLeft method works by applying the function/operation you give it, and applying it to successive elements in the collection

Acknowledgement & References

- Acknowledgement
 - Some slides are adapted from slides by Jure Leskovec et al. <http://www.mmds.org>
- References
 - <http://infolab.stanford.edu/~ullman/mmds/ch11.pdf>
 - <http://www.mmds.org>

*Be Scalable
for both computing & living*